

Designing Studies and Evaluating Research Results: Type M and Type S Errors for Pearson Correlation Coefficient

Giulia Bertoldo ¹, Claudio Zandonella Callegher ¹, and Gianmarco Altoè ^{1,*}

¹ Department of Developmental Psychology and Socialisation, University of Padova,
Padova, Italy

Submitted to *Meta-Psychology*. Click here to follow the fully transparent editorial process of this submission. Participate in open peer review by commenting through hypothes.is directly on this preprint.

Author Note

 ORCID

Giulia Bertoldo: 0000-0002-6960-3980

Claudio Zandonella Callegher: 0000-0001-7721-6318

Gianmarco Altoè: 0000-0003-1154-9528

We have no known conflict of interest to disclose.

*Correspondence:

Gianmarco Altoè, Department of Developmental Psychology and Socialization,
University of Padova, Via Venezia 8, 35131 Padova, Italy
gianmarco.altoe@unipd.it

Abstract

Keywords— Correlation coefficient; Type M Error, Type S error, Design analysis, Effect size

1 Introduction

Psychological science is increasingly committed to scrutinize its published findings by promoting large-scale replication efforts (Camerer et al., 2016, 2018; Ebersole et al., 2016; Klein et al., 2014, 2018; Open Science Collaboration, 2015). In replication studies, the protocol of a previously published study is repeated as closely as possible with a new sample. Interestingly, many replication studies found smaller effects than originals, suggesting a sort of decline effect from original studies to replications (Camerer et al., 2018; Open Science Collaboration, 2015; Schooler, 2014). Among the many factors that could explain such decline, one relates to a feature of study design: statistical power (Protzko & Schooler, 2017). In particular, it is plausible for original studies to have lower statistical power than their replications. When underpowered studies are analyzed using thresholds, such as statistical significance levels, effects passing such thresholds have to exaggerate the true effect size (Button et al., 2013; Gelman, Skardhamar, & Aaltonen, 2017; Ioannidis, 2008; Ioannidis, Pereira, & Horwitz, 2013; Lane & Dunlap, 1978). Indeed, as will be extensively shown below, in underpowered studies only large effects correspond to values that can reject the null hypothesis and be statistically significant. As a consequence, if the original study was underpowered and found an exaggerated estimate of the effect, the replication results will likely show a decline. The concept of statistical power finds its natural development in the Neyman-Pearson framework of statistical inference and this is the framework that we adopt in this contribution. Further discussion on the Neyman and Pearson approach and a comparison with Null Hypothesis Significance Testing (NHST) is available in Altoè et al. (2020) and Gigerenzer, Krauss, and Vitouch (2004).

Effect size exaggeration is a corollary consequence of low statistical power that might not be evident at first. Indeed, by definition, statistical power focuses on statistical significance, not effect size estimation (i.e., statistical power is the probability to find a statistically significant result if an effect of a certain size actually exists). However, effect size estimation and statistical significance are closely related. This point can be highlighted estimating the Type M (magnitude) and Type S (sign) errors characterizing a study design (Gelman & Carlin, 2014). Given a study design (i.e. sample size, statistical test directionality, and α level), Type M error indicates the factor by which a statistically significant effect would be, on average, exaggerated. Type S error indicates the probability to find a statistically significant effect in the opposite

direction to the one considered plausible. It is evident that both errors are defined starting from a reasoned guess on the plausible magnitude and direction of the effect under study, which is called *plausible effect size* (Gelman & Carlin, 2014). The analysis that researchers perform to consider Type M and Type S errors in their research practice is called *design analysis*, given the special focus posed into considering the design of a study (Altoè et al., 2020; Gelman & Carlin, 2014).

Why do these errors matter? The exaggeration of effect sizes, in the right or in the wrong direction, has important consequences on a theoretical and applied level. On a theoretical level, studies' designs with high Type M and Type S errors can foster distorted expectations on the effect under study, triggering a vicious cycle for the planning of future studies. This point is relevant also for the design of replication studies, which could turn out to be underpowered if they do not take into account possible inflations of the original effect (Button et al., 2013). When studies are used to inform policymaking and real-world interventions, implications can go beyond the academic research community and can impact society at large. In these settings, we could assist to a “hype and disappointment cycle” (Gelman, 2019), where true effects turn out to be much less impressive than expected. This can produce undesirable consequences on people's lives, a consideration that invites researchers to assume responsibility in effectively communicating the risks related to effects quantification.

Type M (magnitude) and Type S (sign) errors are not widely known in the psychological research community but their consideration during the research process has the potential to improve current research practices, for example, by increasing the awareness that design choices have on possible studies' results. In a previous work we illustrated Type M and Type S errors using Cohens' d as a measure of effect size (Altoè et al., 2020). In this contribution we aim to further increase familiarity with Type M and Type S errors, considering another common effect size measures in psychology: Pearson correlation coefficient, ρ . In the following paragraphs we discuss what Type M and Type S errors are and how to perform a design analysis using the ad-hoc R functions that we developed.

2 Type M and Type S errors

Pearson correlation coefficient is a standardized effect size measure indicating the strength and the direction of the relationship between two continuous variables (Cohen, 1988; Ellis, 2010). Even though the correlation coefficient is widely known, we briefly go over its main features using an example. Imagine that we were interested to measure the relationship between anxiety and depression in a population and we plan a study with n participants, where, for each participant, we measure the level of anxiety (i.e. variable X) and the level of depression (i.e. variable Y). At the end of the study we will have n pairs of values X and Y . The correlation coefficient helps us answer the questions: how strong is the relationship between anxiety and depression in this population? Is the relationship positive or negative? Correlation ranges from -1

to $+1$, indicating respectively two extreme scenarios of perfect negative relationship and perfect positive relationship. Since the correlation coefficient is a dimensionless number, it is a signal to noise ratio where the signal is given by the covariance between the two variables ($cov(x, y)$) and the noise is expressed by the product between the standard deviations of the two variables ($S_x S_y$; see Formula 1). Following the conventional standards, in this contribution we will use the symbol ρ to indicate the correlation in the population and the symbol r to indicate the value measured in a sample.

$$r = \frac{cov(x, y)}{S_x S_y}. \quad (1)$$

Magnitude and sign are two important features characterizing Pearson correlation coefficient and effect size measures in general. And, when estimating effect sizes, errors could be committed exactly regarding these two aspects. Gelman and Carlin (2014) introduced two indexes to quantify these risks:

- Type M error, where M stands for magnitude, is also called Exaggeration Ratio - the factor by which a statistically significant effect is on average exaggerated.
- Type S error, sign - the probability to find a statistically significant result in the opposite direction to the plausible one.

How are these errors computed? In the next paragraphs we approach this question preferring an intuitive perspective. For a formal definition of these errors, we refer the reader to Altoè et al. (2020); Gelman and Carlin (2014); Lu, Qiu, and Deng (2018). Take as an example the previous fictitious study on the relationship between anxiety and depression and imagine we decide to sample 50 individuals (sample size, $n = 50$) and to set the α level to 5% and to perform a two-tailed test. On the basis of theoretical considerations, we expect the plausibly true correlation in the population to be quite strong and positive which we formalize as $\rho = .50$. In order to evaluate the Type M and Type S errors in this research design, imagine repeating the same study many times with new samples drawn from the same population and, for each study, register the observed correlation (r) and the corresponding p-value.

The first step to compute Type M error is to select only the observed correlation coefficients that are statistically significant in absolute value (for the moment, we do not care about the sign) and to calculate their mean. Type M error is given by the ratio between this mean (i.e., mean of statistically significant correlation coefficients in absolute value) and the plausible effect hypothesized at the beginning, which in this example is $\rho = .50$. Thus, given a study design, Type M error tells us what is the average overestimation of an effect that is statistically significant.

Type S error is computed as the proportion of statistically significant results that have the opposite sign compared to the plausible effect size. In the present example we hypothesized a positive relationship, specifically $\rho = .50$. Then, Type S error is the ratio between the number of times we observed a negative statistically significant result and the total number of statistically significant results. In other words, Type

S error indicates the probability to obtain a statistically significant result in the opposite direction to the one hypothesized.

The central and possibly most difficult point in this process is reasoning on what could be the plausible magnitude and direction of the effect of interest. This critical process, which is central also in traditional a priori power analysis, is an opportunity for researchers to aggregate, formalize and incorporate prior information on the phenomenon under investigation (Gelman & Carlin, 2014). What is plausible can be determined on theoretical grounds, using expert knowledge elicitation techniques and consulting literature reviews and meta-analysis, always taking into account the presence of effect sizes inflation in the published literature (Todo:) (Anderson, 2019). The plausible effect size approximates the true effect, which is never known but can be thought as “that which would be observed in a hypothetical infinitely large sample” (Gelman & Carlin, 2014, p. 642). For a more exhaustive description on plausible effect size, we refer the interested reader to Altoè et al. (2020); Gelman and Carlin (2014).

3 Design Analysis

Researchers can consider Type M and Type S errors in their practice by performing a *design analysis* (Altoè et al., 2020; Gelman & Carlin, 2014). Ideally, a design analysis should be performed when designing a study. In this phase it is specifically called *prospective design analysis* and it can be used as a sample size planning strategy where statistical power is considered together with Type M and Type S errors. However, design analysis can also be beneficial to evaluate the inferential risks in studies that have already been conducted and where the study design is known. In these cases, Type M and Type S errors can support results interpretation by communicating the inferential risks in that research design. When design analysis happens at this later stage, it takes the name of *retrospective design analysis*.

In the following sections we illustrate how to perform prospective and retrospective design analysis using some examples. We developed two R functions to perform design analysis for Pearson correlation, which are available at the page (Todo:) The function to perform a prospective design analysis is called “**pro_r**”. It requires as input the plausible effect size (**rho**), the statistical power (**power**) the directionality of the test (**alternative**) which can be set as: “**two.sided**”, “**less**” or “**greater**” . Type I error rate (**sig_level**) is set as default at 5% and can be changed by the user. The “**pro_r**” function returns the necessary sample size to achieve the desired statistical power, Type M error rate, the Type S error probability, and the critical value(s) above which a statistically significant result can be found. The function to perform retrospective design analysis is called “**retro_r**”. It requires as input the plausible effect size, the sample size used in the study, and the directionality of the test that was performed. Also in this case, Type I error rate is set as default at 5% and can be changed by the user. The function “**retro_r**” returns the Type M error rate, the Type S error probability, and the critical value(s). (Todo:) For further

details on (how to use) computational aspects regarding the R functions, we refer to Appendix A.

4 Case Study

To familiarize the reader with Type M and Type S errors, we start our discussion with a retrospective design analysis of a published study. However, the ideal temporal sequence in the research process would be to perform a prospective design analysis in the planning stage of a research project. This is the time when the design is being laid out and useful improvements can be made to obtain more robust results. In this contribution, the order of presentation aims first, to provide an understanding of how to interpret Type M and Type S errors, and then discuss how they could be minimized. The following case study was chosen for illustrative purposes only and, by no means our objective is to judge the study beyond illustrating an application of how to calculate Type M and Type S errors on a published study.

We consider the study published in *Science* by Eisenberger, Lieberman, and Williams (2003) entitled: “Does Rejection Hurt? An fMRI Study of Social Exclusion”. The research question originated from the observation that the Anterior Cingulate Cortex (ACC) is a region of the brain known to be involved in the experience of physical pain. Could pain from social stimuli, such as social exclusion, share similar neural underpinnings? To test this hypothesis, 13 participants were recruited and each one had to play a virtual game with other two players while undergoing functional Magnetic Resonance Imaging (fMRI). The other two players were fictitious, and participants were actually playing against a computer program. Players had to toss a virtual ball among each other in three conditions: social inclusion, explicit social exclusion and implicit social exclusion. In the social inclusion condition the participant regularly received the ball. In the explicit social exclusion condition the participant was told that, due to technical problems, he was not going to play that round. In the implicit social exclusion condition, the participant experienced being intentionally left out from the game by the other two players. At the end of the experiment, each participant completed a self-report measure regarding their perceived distress when they were intentionally left out by the other players. Considering only the implicit social exclusion condition, a correlation coefficient was estimated between the measure of distress and neural activity in the Anterior Cingulate Cortex. As suggested by the large and statistically significant correlation coefficient between perceived distress and activity in the ACC, (Todo:) $r = .88$, $p < .005$ (Eisenberger et al., 2003, p. 291), authors concluded that social and physical pain seem to share similar neural underpinnings.

Before proceeding to the retrospective design analysis, we refer the interested reader to some background history regarding this study. This was one of the many studies included in the famous paper “Puzzlingly High Correlations in fMRI Studies of Emotion, Personality, and Social Cognition” (Vul, Harris, Winkelman, & Pashler, 2009) which raised important issues regarding the analysis of neuroscientific data. In particular, this paper noted that the magnitude of correlation coefficients between fMRI measures

and behavioral measures were beyond what could be considered plausible. We refer the interested reader also to the commentary by Yarkoni (2009), who noted that the implausibly high correlations in fMRI studies could be largely explained by the low statistical power of experiments.

A retrospective design analysis should start with thorough reasoning on the plausible size and direction of the effect under study. In order to produce valid inferences, a lot of attention should be devoted to this point by integrating external information. For the sake of this example, we turn to the considerations made by Vul and Pashler (2017) who (Todo:) considered correlations between personality measures and neural activity to be likely around $\rho = .25$. A correlation of $\rho = .50$ was deemed plausible but optimistic and a correlation of $\rho = .75$ was considered theoretically plausible but unrealistic.

5 Retrospective Design Analysis

To perform a retrospective design analysis on the case study, we need information on the research design and the plausible effect size. Based on the previous considerations, we set the plausible effect size to be $\rho = .25$. Information on the sample size was not available in the original study (Eisenberger et al., 2003) and was retrieved from Vul et al. (2009) to be $n = 13$. The α level and the directionality of the test were not reported in the original study, so for the purpose of this example we will consider $\alpha = .05$ and a two-tailed test. Given this study design, what are the inferential risks in terms of effect size estimation? We can use the R function “`retro_r`”, which inputs and outputs are displayed (Todo:) below.

```
retro_r(rho = 0.25, n = 13, alternative = "two.sided", sig_level = 0.05, seed = 2020)

##
##  Design Analysis
##
## Hypothesized effect: rho = 0.25
##
## Study characteristics:
##      n      alternative  sig_level
##    13      two.sided      0.05
##
## Inferential risks:
##      power  typeM  typeS
##    0.127   2.583   0.028
##
## Minimum significant r = ±0.553
```

In this study, the statistical power is .13, that is to say, there is a 13% probability to reject the null hypothesis, if an effect of at least $\rho = |.25|$ exists. Consider this point together with the results obtained in the experiment: $r = .88$, $p < .005$ (Eisenberger et al., 2003, p. 291). It is clear that, even though the probability to reject the null hypothesis is low (power of 13%), this event could happen. And when

it does happen, it is tempting to believe that results are (Todo:) impressive (Gelman & Loken, 2014). However, this design comes with serious inferential risks for the estimation of effect sizes, which could be grasped by presenting Type M and Type S errors. A glance at their value communicates that it is not impossible to find a statistically significant result, but when it does happen, the effect sizes could be largely overestimated - Type M = 2.58 - and maybe even in the wrong direction - Type S = .03. The Type M error rate of (Todo:) around 2.60 indicates that a statistically significant correlation is on average about two and a half times the plausible value. In other words, statistically significant results emerging in such a research design will on average overestimate the plausible correlation coefficient by 160%. The Type S error of .03 suggests that there is a three percent probability to find a statistically significant result in the opposite direction, in this example, a negative relationship.

In this research design, the critical values above which a statistically significant result is declared correspond to $\rho = \pm .55$ ((Todo:) see output above). These values are highlighted in Figure 1 as the vertical lines in the sampling distribution of correlation coefficients under the null hypothesis. Notice that the plausible effect size lies in the region of acceptance of the null hypothesis. Therefore, it is impossible to simultaneously find a statistically significant result and estimate an effect close to the plausible one ($\rho = .25$). The figure represents the so called Winner's curse: "the 'lucky' scientist who makes a discovery is cursed by finding an inflated estimate of that effect" (Button et al., 2013).




Figure 1: Winner's course. H_0 = Null Hypothesis, H_1 = Alternative Hypothesis. When sample size, directionality of the test and Type error probability are set, also the smallest effect size above which is possible to find a statistically significant result is set. In this case, the plausible effect size, $\rho = .25$, lies in the region where it is not possible to reject H_0 . Thus, it is impossible to simultaneously find a statistically significant result and an effect close to the plausible one. In other words, a statistically significant effect must exaggerate the plausible effect size.

6 Prospective Design Analysis

Ideally, Type M and Type S errors should be considered in the design phase of a study during the decision-making process regarding the experimental protocol. At this stage, prospective design analysis can be used as a sample size planning strategy which aims to minimize Type M and Type S errors in the upcoming study.

Imagine that we were part of the research team in the previous case study exploring the relationship between activity in the Anterior Cerebral Cortex and perceived distress. When drafting the research protocol, we face the inevitable discussion on how many participants we are going to recruit. This choice depends on available resources, type of study design, constraints of various nature and, importantly, the plausible magnitude and direction of the phenomenon that we are going to study. As previously mention, deciding on a plausible effect size is a fundamental step which requires great effort and should not be (Todo:) a game where different numbers are tried until the desired sample size is reached. Instead, proposing a plausible effect size is where the expert knowledge of the researcher can be formalized and can greatly contribute to the informativeness of the study that is being planned. For the sake of this examples, we adopt the previous consideration and we suppose that common agreement is reached on a plausible correlation coefficient to be around $\rho = .25$. Finally, we would like to leave open the possibility to explore whether the relationship goes in the opposite direction to the one hypothesized, so we decide to perform a two-tailed test.

We can implement the prospective design analysis using the function “`pro_r`” which inputs and outputs are displayed in (Todo:) Figure 3. About 125 participants are necessary to have 80 % probability to detect an effect of at least $\rho = \pm .25$, if it actually exists. With this sample size, the Type S error is minimized and approximates zero. In this study design, the Type M error is 1.11 indicating that statistically significant results are on average exaggerated by 11%. It is possible to notice that the critical values are $\rho = \pm .18$, further highlighting that our plausible effect size is actually included among those values that lead to the acceptance of the alternative hypothesis.

```
pro_r(rho = 0.25, power = 0.8, alternative = "two.sided", sig_level = 0.05, seed = 2020)

##
## Design Analysis
##
## Hypothesized effect: rho = 0.25
##
## Study characteristics:
##   n      alternative  sig_level
##  122    two.sided     0.05
##
## Inferential risks:
##   power  typeM  typeS
```

```
##      0.8      1.118    0
##
## Minimum significant r = ±0.178
```

In a design analysis, it is advisable to investigate how the inferential risks would change according to different scenarios in terms of statistical power and plausible effect size. Changes in both these factors impact Type M and Type S errors. For example, maintaining the plausible correlation of $\rho = .25$, if we decrease statistical power from .80 to .60 only 77 participants are required (see Table 1). However, this is associated with an increased Type M error rate from 1.12 to 1.28. That is to say, with 77 subjects the plausible effect size will be on average overestimated by 28%. Alternatively, imagine that we would like to maintain a statistical power of 80%, what happens if the plausible effect size is slightly larger or smaller? The necessary sample size would spike to 344 for a $\rho = .15$ and decrease to 60 for $\rho = .35$. In both scenarios, the Type M error remains about 1.12, which reflects the more general point that for 80% power, Type M error is around 1.10. In all these scenario, Type S error is close to zero, hence not worrisome.

Table 1: Prospective design analysis in different scenarios of plausible effect size and statistical power.

| ρ | Power | Sample Size | Type M | Type S | Minimum sign. r |
|--------|-------|-------------|--------|--------|-------------------|
| 0.25 | 0.6 | 76 | 1.280 | 0 | ± 0.226 |
| 0.15 | 0.8 | 344 | 1.116 | 0 | ± 0.106 |
| 0.35 | 0.8 | 60 | 1.115 | 0 | ± 0.254 |

For completeness, Figure ?? summarizes the relationship between statistical power, Type M and Type S errors as a function of sample size in three scenarios of plausible correlation coefficients. We display the three values that Vul and Pashler (2017) considered for correlations between fMRI measures and behavioral measures with different degrees of plausibility. An effect of $\rho = .75$ was deemed theoretically plausible but unrealistic, $\rho = .50$ was more plausible but optimistic, and $\rho = .25$ was more likely. The curves illustrate a general point: Type M and Type S error increase with smaller sample sizes, smaller plausible effect sizes and lower statistical power. Also, the figure shows that statistical power, Type M and Type S errors are related to each other: as power increases, Type M and Type S errors decrease.

At first, it might seem that Type M and Type S errors are redundant with the information provided by statistical power. Even though they are related, we believe that Type M and Type S errors bring added value during the design phase of a research protocol because they facilitate a connection between how a study is planned and how results will actually be evaluated. That is to say, final results will comprise of a test statistics with an associated p-value and effect size measure. If the interest is maximizing the accuracy with which effects will be estimated, then Type M and Type S errors directly communicate the consequences of design choices on effect size estimation.

7 Session Information

```
sessionInfo(package = NULL)

## R version 3.6.1 (2019-07-05)
## Platform: x86_64-apple-darwin15.6.0 (64-bit)
## Running under: macOS Mojave 10.14.6
##
## Matrix products: default
## BLAS: /Library/Frameworks/R.framework/Versions/3.6/Resources/lib/libRblas.0.dylib
## LAPACK: /Library/Frameworks/R.framework/Versions/3.6/Resources/lib/libRlapack.dylib
##
## locale:
## [1] en_US.UTF-8/en_US.UTF-8/en_US.UTF-8/C/en_US.UTF-8/en_US.UTF-8
##
## attached base packages:
## [1] stats      graphics  grDevices  utils      datasets  methods   base
##
## other attached packages:
## [1] docstring_1.0.0 MASS_7.3-51.6 forcats_0.5.0 stringr_1.4.0 dplyr_0.8.5
## [6] purrr_0.3.4 readr_1.3.1 tidyr_1.0.2 tibble_3.0.1 ggplot2_3.3.0
## [11] tidyverse_1.3.0 knitr_1.28
##
## loaded via a namespace (and not attached):
## [1] tidyselect_1.0.0 xfun_0.12 haven_2.2.0 lattice_0.20-38
## [5] colorspace_1.4-1 vctrs_0.2.4 generics_0.0.2 rlang_0.4.5
## [9] pillar_1.4.3 glue_1.3.1 withr_2.1.2 DBI_1.1.0
## [13] dbplyr_1.4.3 tikzDevice_0.12.3 modelr_0.1.6 readxl_1.3.1
## [17] lifecycle_0.2.0 munsell_0.5.0 gtable_0.3.0 cellranger_1.1.0
## [21] rvest_0.3.5 codetools_0.2-16 evaluate_0.14 fansi_0.4.1
## [25] highr_0.8 broom_0.5.6 Rcpp_1.0.3 formatR_1.7
## [29] scales_1.1.0 backports_1.1.5 filehash_2.4-2 jsonlite_1.6.1
## [33] fs_1.4.1 digest_0.6.25 hms_0.5.3 stringi_1.4.6
## [37] grid_3.6.1 cli_2.0.2 tools_3.6.1 magrittr_1.5
## [41] crayon_1.3.4 pkgconfig_2.0.3 ellipsis_0.3.0 xml2_1.2.2
## [45] reprex_0.3.0 lubridate_1.7.8 assertthat_0.2.1 httr_1.4.1
## [49] roxygen2_7.1.0 rstudioapi_0.11 R6_2.4.1 nlme_3.1-141
## [53] compiler_3.6.1
```

Appendix A

References

- Altoè, G., Bertoldo, G., Zandonella Callegher, C., Toffalini, E., Calcagni, A., Finos, L., & Pastore, M. (2020, January). Enhancing Statistical Inference in Psychological Research via Prospective and Retrospective Design Analysis. *Frontiers in Psychology*, 10. doi: <https://doi.org/10.3389/fpsyg.2019.02893>
- Button, K., Ioannidis, J., Mokrysz, C., Nosek, B., Flint, J., Robinson, E., & Munafò, M. (2013). Power failure: Why small sample size undermines the reliability of neuroscience. *Nature Reviews Neuroscience*, 14(5), 365–376. doi: <https://doi.org/10.1038/nrn3475>
- Camerer, C. F., Dreber, A., Forsell, E., Ho, T.-H., Huber, J., Johannesson, M., ... Wu, H. (2016). Evaluating replicability of laboratory experiments in economics. *Science*, 351(6280), 1433–1436. doi: <https://doi.org/10.1126/science.aaf0918>
- Camerer, C. F., Dreber, A., Holzmeister, F., Ho, T.-H., Huber, J., Johannesson, M., ... Wu, H. (2018, September). Evaluating the replicability of social science experiments in Nature and Science between 2010 and 2015. *Nature Human Behaviour*, 2(9), 637–644. doi: <https://doi.org/10.1038/s41562-018-0399-z>
- Cohen, J. (1988). *Statistical power analysis for the behavioral sciences*. Lawrence Erlbaum Associates.
- Ebersole, C. R., Atherton, O. E., Belanger, A. L., Skulborstad, H. M., Allen, J. M., Banks, J. B., ... Nosek, B. A. (2016). Many Labs 3: Evaluating participant pool quality across the academic semester via replication. *Journal of Experimental Social Psychology*, 67, 68–82. doi: <https://doi.org/10.1016/j.jesp.2015.10.012>
- Eisenberger, N. I., Lieberman, M. D., & Williams, K. D. (2003). Does rejection hurt? An fMRI study of social exclusion. *Science*, 302(5643), 290–292. doi: <https://doi.org/10.1126/science.1089134>
- Ellis, P. D. (2010). *The Essential Guide to Effect Sizes*. Cambridge University Press.
- Gelman, A. (2019). *Embracing Variation and Accepting Uncertainty: Implications for Science and Metascience* [Video]. (Metascience 2019 Symposium: The Emerging Field of Research on the scientific Process. September 5th–8th, Stanford University. <https://www.metascience2019.org/presentations/andrew-gelman/>)
- Gelman, A., & Carlin, J. (2014, November). Beyond Power Calculations: Assessing Type S (Sign) and Type M (Magnitude) Errors. *Perspectives on Psychological Science*, 9(6), 641–651. doi: <https://doi.org/10.1177/1745691614551642>
- Gelman, A., & Loken, E. (2014). The statistical crisis in science. *American scientist*, 102(6), 460–466. doi: <https://doi.org/10.1511/2014.111.460>
- Gelman, A., Skardhamar, T., & Aaltonen, M. (2017). Type M Error Might Explain Weisburd’s Paradox. *Journal of Quantitative Criminology*. doi: <https://doi.org/10.1007/s10940-017-9374-5>
- Gigerenzer, G., Krauss, S., & Vitouch, O. (2004). The Null Ritual: What You Always Wanted to Know About Significance Testing but Were Afraid to Ask. In *The SAGE Handbook of Quantitative Methodology for the Social Sciences* (pp. 392–409). 2455 Teller Road, Thousand Oaks California 91320 United States of America: SAGE Publications, Inc. doi: <https://doi.org/10.4135/9781412986311.n21>
- Ioannidis, J. P. A. (2008, September). Why Most Discovered True Associations Are Inflated:. *Epidemiology*, 19(5), 640–648. doi: <https://doi.org/10.1097/EDE.0b013e31818131e7>
- Ioannidis, J. P. A., Pereira, T. V., & Horwitz, R. I. (2013, February). Emergence of Large Treatment Effects From Small Trials—Reply. *JAMA*, 309(8), 768–769. doi: <https://doi.org/10.1001/jama.2012.208831>
- Klein, R. A., Ratliff, K. A., Vianello, M., Adams, R. B., Bahník, Š., Bernstein, M. J., ... Nosek, B. A. (2014, January). Investigating Variation in Replicability. *Social Psychology*, 45(3), 142–152. doi: <https://doi.org/10.1027/1864-9335/a000178>

- Klein, R. A., Vianello, M., Hasselman, F., Adams, B. G., Reginald B. Adams, J., Alper, S., ... Nosek, B. A. (2018). Many labs 2: Investigating variation in replicability across samples and settings. *Advances in Methods and Practices in Psychological Science*, 1(4), 443–490. doi: <https://doi.org/10.1177/2515245918810225>
- Lane, D. M., & Dunlap, W. P. (1978). Estimating effect size: Bias resulting from the significance criterion in editorial decisions. *British Journal of Mathematical and Statistical Psychology*, 31(2), 107–112. doi: <https://doi.org/10.1111/j.2044-8317.1978.tb00578.x>
- Lu, J., Qiu, Y., & Deng, A. (2018, March). A note on Type S/M errors in hypothesis testing. *British Journal of Mathematical and Statistical Psychology*. doi: <https://doi.org/10.1111/bmsp.12132>
- Open Science Collaboration. (2015, August). Estimating the reproducibility of psychological science. *Science*, 349(6251), aac4716–aac4716. doi: <https://doi.org/10.1126/science.aac4716>
- Protzko, J., & Schooler, J. W. (2017). Decline effects: Types, mechanisms, and personal reflections. In *Psychological science under scrutiny: Recent challenges and proposed solutions*. (pp. 85–107). Wiley-Blackwell. doi: <https://doi.org/10.1002/9781119095910.ch6>
- Schooler, J. (2014). Turning the Lens of Science on Itself: Verbal Overshadowing, Replication, and Metascience. *Perspectives on Psychological Science*, 9(5), 579–584. doi: <https://doi.org/10.1177/1745691614547878>
- Vul, E., Harris, C., Winkielman, P., & Pashler, H. (2009). Puzzlingly high correlations in fMRI studies of emotion, personality, and social cognition. *Perspectives on Psychological Science*, 4(3), 274–290. doi: <https://doi.org/10.1111/j.1745-6924.2009.01125.x>
- Vul, E., & Pashler, H. (2017). Suspiciously high correlations in brain imaging research. In *Psychological science under scrutiny* (pp. 196–220). John Wiley & Sons, Ltd. doi: <https://doi.org/10.1002/9781119095910.ch11>
- Yarkoni, T. (2009, May). Big Correlations in Little Studies: Inflated fMRI Correlations Reflect Low Statistical Power—Commentary on Vul et al. (2009). *Perspectives on Psychological Science*, 4(3), 294–298. doi: <https://doi.org/10.1111/j.1745-6924.2009.01127.x>